



Dialogue Breakdown Detection in Chatbots

student: Mariya Hendriksen
supervisor: Prof. Dr. Marie-Francine Moens
assessor: Prof. Dr. Hugo Van Hamme
mentor: Artuur Leeuwenberg

Faculty of Engineering Science

February 1, 2019

Outline

Introduction

Data & Methods

- Dataset

- Model Type & Architecture

- Word Embedding Models

Implementation

- Data Preprocessing

- Models Development

Evaluation

- Evaluation Metrics

- Evaluation Results

Conclusion

Outline

Introduction

Data & Methods

- Dataset

- Model Type & Architecture

- Word Embedding Models

Implementation

- Data Preprocessing

- Models Development

Evaluation

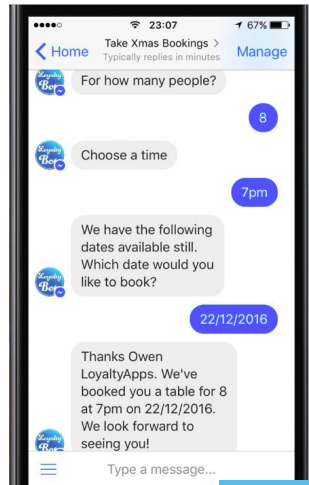
- Evaluation Metrics

- Evaluation Results

Conclusion

Chatbot

- A chatbot is a dialogue system which conducts a conversation with a user via text or speech.
- The popularity of chatbots increase.
- Instant 24 hours service
- Usage categories: website help, education, customer service, e-commerce, entertainment, etc.



Dialogue Breakdown

- Breakdown is a point in a dialogue when the interaction is interrupted with or without completion of the performed task [Martinovsky and Traum, 2006]

It's nice to *go shopping alone*.

I agree. That's nice.

Shopping takes time.

Window shopping is also fun.

It's fun to go shopping with somebody.

Chatbot utterances are red.

About the Project

- Aim: create a dialogue breakdown detection model
- Research questions:
 - Before building the model
 - What models for breakdown detection are available?
 - Which types of models perform better?
 - After building the model
 - How does the architecture of the model influence its performance?
 - How does the embedding type of the detector influence its performance?

Contributions

- Overview and comparison of the existing dialogue breakdown detectors
- Introduction of dialogue breakdown detection models which outperform the baseline
- Exploration of relationship between model architecture and its performance
- Investigation of the relationship between embedding type and detector's performance

Related Literature

- Dialogue breakdown detection challenge:
 - introduced in [Ryuichiro et al., 2015]
 - three challenges held [Higashinaka et al., 2016, 2017]
 - fourth challenge will be held during IWSDS 2019 [dbd]
- Applications:
 - re-ranking system responses [Inaba and Takahashi]
 - aggregating utterances that can be used as system responses [Sugiyama, 2016]

Outline

Introduction

Data & Methods

- Dataset

- Model Type & Architecture

- Word Embedding Models

Implementation

- Data Preprocessing

- Models Development

Evaluation

- Evaluation Metrics

- Evaluation Results

Conclusion

Dataset

- Dataset from the Dialogue Breakdown Detection Challenge 3 [dbdc33]
- Dialogues sources:
 - chatbots IRIS, Tick-Tock, 'Yura and Idris'
 - Conversational Intelligence Challenge
- Labels: NB, PB, B
- Classes are balanced
- 30 annotators per system utterance

Baseline Definition

Based on the overview of existing models performance, we chose two models for the baseline:

- accuracy baseline
- $F_1(B)$ and $F_1(PB + B)$ baseline

Baseline type	Model	Score
Accuracy	LSTM + word2vec	0.44
$F_1(B)$	MemNN + attention	0.36
$F_1(PB + B)$	MemNN + attention	0.87

Table: Baseline Models

Model Type & Architecture

- Model Type: Why LSTM?
 - can process sequential data
 - handles long-term dependencies
- Selected architecture types:
 - vanilla LSTM
 - stacked LSTM
 - bidirectional LSTM

Word Embedding Models Selection

- Limited data \implies pretrained word embedding models
- For the experiments, we select three word embedding types:
 - word2vec pretrained on Google News, 300D
 - GloVe pretrained on Twitter, 200D
 - GloVe pretrained on Common Crawl, 300D

Outline

Introduction

Data & Methods

- Dataset

- Model Type & Architecture

- Word Embedding Models

Implementation

- Data Preprocessing

- Models Development

Evaluation

- Evaluation Metrics

- Evaluation Results


Conclusion

Datasets Preprocessing

- Preprocessing steps:
 - General steps:
 - tokenization (with `nltk.tokenize.casual.TweetTokenizer`)
 - apostrophes contractions replacement, e.g., *that's* → *that is*
 - Additional steps:
 - for word2vec pretrained on Google News: punctuation signs removal
 - for GloVe pretrained on tweets: lowercasing
- Dialogues length varies \implies padding

Representations

- **Token representation** obtained with pretrained embedding model
- **Utterance representation** is the average of the token embeddings comprising the utterances
- **Dialogue representation** is a 2D tensor comprising user and system utterances
- Each **label** is represented with one-hot encoding



Models Development I

- **Software:** Python 3.6 with libraries NumPy, NLTK, gensim, matplotlib, sklearn, json. Tensorflow(backend) + Keras
- **Layers:** one input layer, two hidden layers, one output layer
- **Loss function:** categorical cross-entropy:

$$H(y, \hat{y}) = - \sum_{i=1} y_i \log(\hat{y}_i) \quad (1)$$

where y - ground truth label, \hat{y} - predicted label

Models Development II

- **Mini-batching** with batch size 32
- **Optimization method**: Root Mean Square Propagation (RMSProp)
- **Maximum number of epochs**: 100
- To prevent **overfitting**:
 - Early stopping with patience 5
 - Dropout and recurrent dropout, coefficients set to 0.1

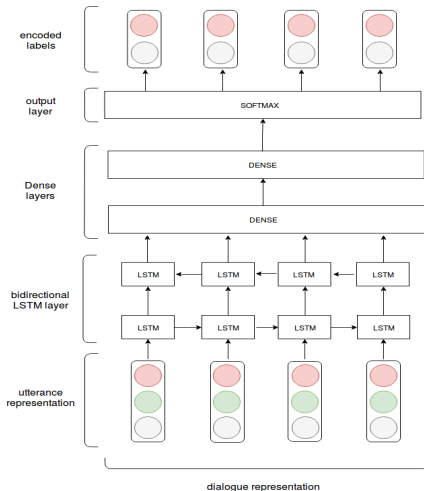
Bidirectional LSTM

```
# imports
from keras.models import Sequential
from keras.layers import LSTM, Dense, Bidirectional

# model definition
bi_lstm = Sequential()
# input layer
lstm1 = Bidirectional(LSTM(NUM_CELLS, return_sequences=True, dropout
    ↳ = DROPOUT_COEFF, recurrent_dropout = DROPOUT_COEFF),
    ↳ input_shape=(MAX_UTTERANCE_LENGTH, UTTERANCE_DIMENSIONALITY),
    ↳ merge_mode = 'sum')
bi_lstm.add(lstm1)
# hidden layers
bi_lstm.add(add(Dense(NUM_CELLS, activation = 'relu')))
bi_lstm.add(Dense(NUM_CELLS, activation = 'relu'))
# output layer
bi_lstm = Dense(NB_LABELS, activation = 'softmax')
bi_lstm.add(outputs)

# model compilation
bi_lstm.compile(
    optimizer = 'rmsprop',
    loss = 'categorical_crossentropy',
    metrics=['accuracy'])

# model training
bi_lstm_history = bi_lstm.fit(
    X_train, y_train,
    batch_size=BATCH_SIZE,
    validation_data = (X_test, y_test),
    callbacks = [EarlyStopping(monitor='val_loss',
    ↳ patience=PATIENCE)],
    epochs=NUM_EPOCHS)
```



Outline

Introduction

Data & Methods

- Dataset

- Model Type & Architecture

- Word Embedding Models

Implementation

- Data Preprocessing

- Models Development

Evaluation

- Evaluation Metrics

- Evaluation Results

Conclusion

Evaluation Metrics

- **Accuracy:** what fraction of items is classified correctly?
- **F_1 score:**
 - Incorporates precision and recall:
 - Precision: how many selected items are relevant?
 - Recall: how many relevant items are selected?
 - A variant of F-score with precision and recall equally important

Models Performance

Model	Embedding	Accuracy	$F_1(B)$	$F_1(PB + B)$
Vanilla LSTM	Google News	0.42	0.38	0.84
	GloVe Twitter	0.41	0.31	0.78
	GloVe Common Crawl	0.44	0.46	0.93
Stacked LSTM	Google News	0.39	0.30	0.77
	GloVe Twitter	0.42	0.41	0.82
	GloVe Common Crawl	0.42	0.45	0.93
Bi-LSTM	Google News	0.42	0.33	0.75
	GloVe Twitter	0.44	0.34	0.83
	GloVe Common Crawl	0.46	0.37	0.85

Table: Models performance results. The best performance is indicated in bold

Comparison with the Baseline

Model type	Description	Accuracy	$F_1(B)$	$F_1(PB + B)$
Accuracy baseline	LSTM + word2vec Google News	<u>0.44</u>	0.29	0.74
Best accuracy score model	Bi-LSTM + GloVe Common Crawl	<u>0.46</u>	0.37	0.85
F_1 baseline	MemNN + attention	0.29	<u>0.36</u>	<u>0.87</u>
Best F_1 score model	Vanilla LSTM + GloVe Common Crawl	0.44	<u>0.46</u>	<u>0.93</u>

Table: Comparison of the created models with the baseline models.
The best performance is indicated in bold

Relationship Exploration

- LSTM architecture in terms of model performance:
 1. Vanilla LSTM
 2. Stacked LSTM
 3. Bidirectional LSTM
- Word embedding model in terms of detector performance:
 1. GloVe Common Crawl
 2. GloVe Twitter
 3. word2vec Google News

Patterns in Error Types

As the number of unknown tokens increase, the model performance decrease.

- Emoticons: simple (😄 😞) vs. complex (😄 😞)
- Apostrophe contractions:
 - standard, e.g., *isn't* → *is not*
 - ambiguous, e.g., *he's* → *he has* or *he is*
- Misspellings:
 - standard, e.g., 'see', 'tomorrow'
 - merging words, e.g., 'questionWho', 'wait.Where'
- Abbreviations, e.g., *ConvAI* (Conversational Artificial Intelligence)
- Extra:
 - meaningless sequences of characters (e.g., *MAMAXMAMAX*)
 - tokens representing words from a different language

Outline

Introduction

Data & Methods

- Dataset

- Model Type & Architecture

- Word Embedding Models

Implementation

- Data Preprocessing

- Models Development

Evaluation

- Evaluation Metrics

- Evaluation Results

Conclusion

Summary

First part:

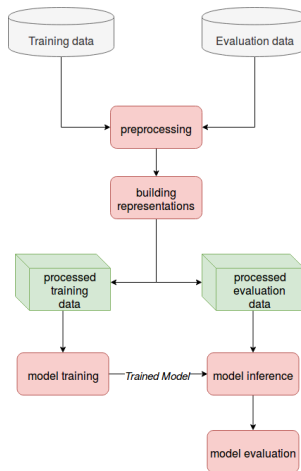
- dataset exploration
- existing models overview
- selection of:
 - baseline
 - model for experiments

Second part:

- dataset preprocessing
- model development

Third part:

- models evaluation
- comparison with the baseline
- exploration of error patterns



Future Work

- Model architecture: investigate further architecture types, e.g., LSTM with attention
- Word embedding model: experiment with Embeddings from Language Model (ELMo)
- Dataset expansion: including different languages and modalities
- Working with the unknown tokens
- Submitting the model for dialogue breakdown detection challenge 4

Thank you for your attention!

Bibliography I

Dialogue breakdown detection challenge 4. URL: <https://sites.google.com/site/dialoguebreakdowndetection4/>, last checked on 30-01-2019.

dbdc33. Dataset. URL: <https://dbd-challenge.github.io/dbdc3/datasets>, last checked on 24-01-2019.

Ryuichiro Higashinaka, Kotaro Funakoshi, Michimasa Inaba, Yuki Arase, and Yuiko Tsunomori. The dialogue breakdown detection challenge 2. *Proceedings of SIG-SLUD*, 2016.

Bibliography II

Ryuichiro Higashinaka, Kotaro Funakoshi, Michimasa Inaba, Yuiko Tsunomori, Tetsuro Takahashi, and Nobuhiro Kaji. Overview of dialogue breakdown detection challenge 3. *Proceedings of Dialog System Technology Challenge*, 6, 2017.

Michimasa Inaba and Kenichi Takahashi. Improving the performance of chat-oriented dialogue systems via dialogue breakdown detection.

José Lopes. How generic can dialogue breakdown detection be? the kth entry to dbdc3.

Bilyana Martinovsky and David Traum. The error is the clue: Breakdown in human-machine interaction. Technical report, Institute for Creative Technologies, University of Southern California, 2006.

Bibliography III

Higashinaka Ryuichiro, Kotaro Funakoshi, Yuka Kobayashi, and Inaba Michimasa. The dialogue breakdown detection challenge: Task description, datasets, and evaluation metrics. 2015.

Hiroaki Sugiyama. Utterance selection based on sentence similarities and dialogue breakdown detection on ntcir-12 stc task. In *NTCIR*, 2016.

Appendix

Task Setting

Dialogue Context:

S1: It's nice to go shopping alone.

U1: I agree. That's nice.

S2: Shopping takes time.

U2: Window shopping is also fun.

Succeeding system utterance:

S3: It's fun to go shopping with somebody.


Input



Dialogue breakdown detector



Single dialogue
breakdown label



Probability distribution
of breakdown labels



Output

Figure: Task setting [Higashinaka et al., 2017]

Dialogue Representation in JSON

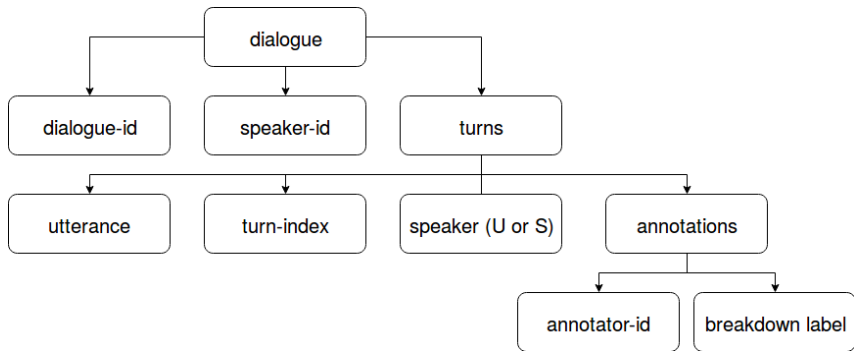


Figure: Dialogue representation in JSON: fields hierarchy

The Dataset Statistics

	Development data				Evaluation data			
	TKTK	IRIS	CIC	YI	TKTK	IRIS	CIC	YI
Dialogues	100	100	115	100	50	50	50	50
Annotators	CF	CF	AMT	AMT	CF	CF	AMT	AMT
NB	35.1%	32.9%	28.9%	34.8%	44.3%	34.5%	29.1%	35.4%
PB	27.6%	27.8%	29.8%	36.1%	29.2%	29.3%	39.3%	40.3%
B	37.3%	39.4%	41.3%	29.1%	26.5%	36.2%	31.6%	24.3%
Fleiss' κ	0.14	0.11	0.05	0.01	0.13	0.09	0.001	-0.006

Table: The dataset statistics

Types of Models

- Conditional Random Fields (CRF)
- Extremely Randomized Trees (ETR)
- Maximum Entropy model (MaxEnt)
- Support Vector Machines (SVM)
- Memory Networks (MemNN)
- Recurrent Neural Networks (RNN). Includes Long Short-term Memory Network (LSTM)
- Models with attention

Model type	Implementation description	Accuracy	$F_1(B)$	$F_1(PB + B)$
CRF	baseline model for DBDC3	0.420	0.354	0.762
MaxEnt	*MaxEnt, labelling exchanges	0.410	0.240	0.220
SVM	SVM + SpeDial feature set	0.340	0.350	0.840
MemNN	MemNN + attention	0.295	0.364	0.874
	MemNN + attention; trained on multilingual data	0.290	0.356	0.870
ETR	ETR + geometric mean	0.426	0.312	0.832
	ETR + cosine similarity of all word pairs	0.431	0.320	0.840
	ETR + arithmetic mean	0.420	0.302	0.835

Model type	Implementation description	Accuracy	$F_1(B)$	$F_1(PB + B)$
RNN	RNN + attention between sentences, GloVe Twitter	0.360	0.208	0.346
	RNN + attention + GloVe Twitter + finetuning	0.210	0.210	0.340
	RNN + attention + GloVe Twitter + extra linguistic features	0.356	0.320	0.805
	LSTM + BoW, word embeddings	0.440	0.290	0.744
	LSTM + BoW, document embeddings	0.422	0.340	0.759
	Hie-Bi-LSTM + GloVe Wikipedia	0.429	0.321	0.763

Table: Models overview: classification results. The best performance is indicated in bold. An asterisk signifies submission without technical paper

Datasets Preprocessing

Decisions to make before preprocessing:

- Keep user utterances? Yes, it improves model performance [Lopes]
- How to feed user utterances? As separate utterances marked with U label

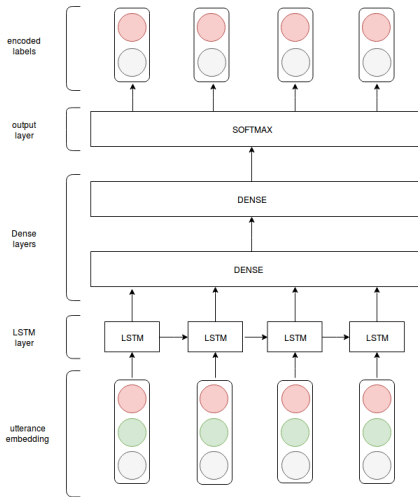
Vanilla LSTM

```
# imports
from keras.models import Sequential
from keras.layers import LSTM, Dense

# model definition
vanilla = Sequential()
# input layer
lstm1 = LSTM(NUM_CELLS, input_shape=(MAX_UTTERANCE_LENGTH,
    ↳ UTTERANCE_DIMENSIONALITY), return_sequences=True, dropout =
    ↳ DROPOUT_COEFF, recurrent_dropout = DROPOUT_COEFF)
vanilla.add(lstm1)
# hidden layers
vanilla.add(Dense(NUM_CELLS, activation = 'relu'))
vanilla.add(Dense(NUM_CELLS, activation = 'relu'))
# output layer
outputs = Dense(NB_LABELS, activation = 'softmax')
vanilla.add(outputs)

# model compilation
vanilla.compile(
    optimizer = 'rmsprop',
    loss = 'categorical_crossentropy',
    metrics=['accuracy'])

# model training
vanilla_history = vanilla.fit(
    X_train, y_train,
    batch_size=BATCH_SIZE,
    validation_data = (X_test, y_test),
    callbacks = [EarlyStopping(monitor='val_loss', patience=5)],
    epochs=NUM_EPOCHS)
```



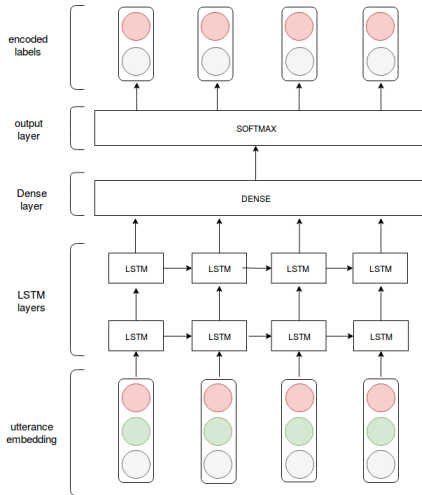
Stacked LSTM

```
# imports
from keras.models import Sequential
from keras.layers import LSTM, Dense

# model definition
stacked = Sequential()
# input layer
lstm1 = LSTM(NUM_CELLS, input_shape=(MAX_UTTERANCE_LENGTH,
    ↳ UTTERANCE_DIMENSIONALITY), return_sequences=True, dropout =
    ↳ DROPOUT_COEFF, recurrent_dropout = DROPOUT_COEFF)
stacked.add(lstm1)
# hidden layers
lstm2 = LSTM(NUM_CELLS, input_shape=(MAX_UTTERANCE_LENGTH,
    ↳ UTTERANCE_DIMENSIONALITY), return_sequences=True, dropout =
    ↳ DROPOUT_COEFF, recurrent_dropout = DROPOUT_COEFF)
stacked.add(lstm2)
vanilla.add(Dense(NUM_CELLS, activation = 'relu'))
# output layer
stacked = Dense(NB_LABELS, activation = 'softmax')
stacked.add(outputs)

# model compilation
stacked.compile(
    optimizer = 'rmsprop',
    loss = 'categorical_crossentropy',
    metrics=['accuracy'])

# model training
stacked_history = stacked.fit(
    X_train, y_train,
    batch_size=BATCH_SIZE,
    validation_data = (X_test, y_test),
    callbacks = [EarlyStopping(monitor='val_loss',
    ↳ patience=PATIENCE)],
    epochs=NUM_EPOCHS)
```



Relationship Exploration

LSTM type	Accuracy	$F_1(B)$	$F_1(PB + B)$
Vanilla LSTM	0.42	0.38	0.85
Stacked LSTM	0.41	0.39	0.84
Bi-LSTM	0.44	0.35	0.81

Table: Average metric scores for every LSTM type

Embedding type	Accuracy	$F_1(B)$	$F_1(PB + B)$
word2vec Google News	0.41	0.34	0.78
GloVe Twitter	0.42	0.35	0.81
GloVe Common Crawl	0.44	0.43	0.90

Table: Average metric scores for every embedding type

Answers to Research Questions

- There are several major types of dialogue breakdown detectors: CRF, ETR, MaxEnt, SVM, MemNN, RNN, and models with attention
- Comparison of the existing models allow concluding that LSTM and MemNN appear to produce better results
- Both model architecture and word embedding model influence the performance. Word embedding model appear to produce a more significant impact.